

# Files and Streams



# Streams

- Most file-related classes are in the java.io package (there is also a java.nio which stands for “Network IO”). Two types of streams, Character or Byte, which can also be buffered or unbuffered.

- Character streams

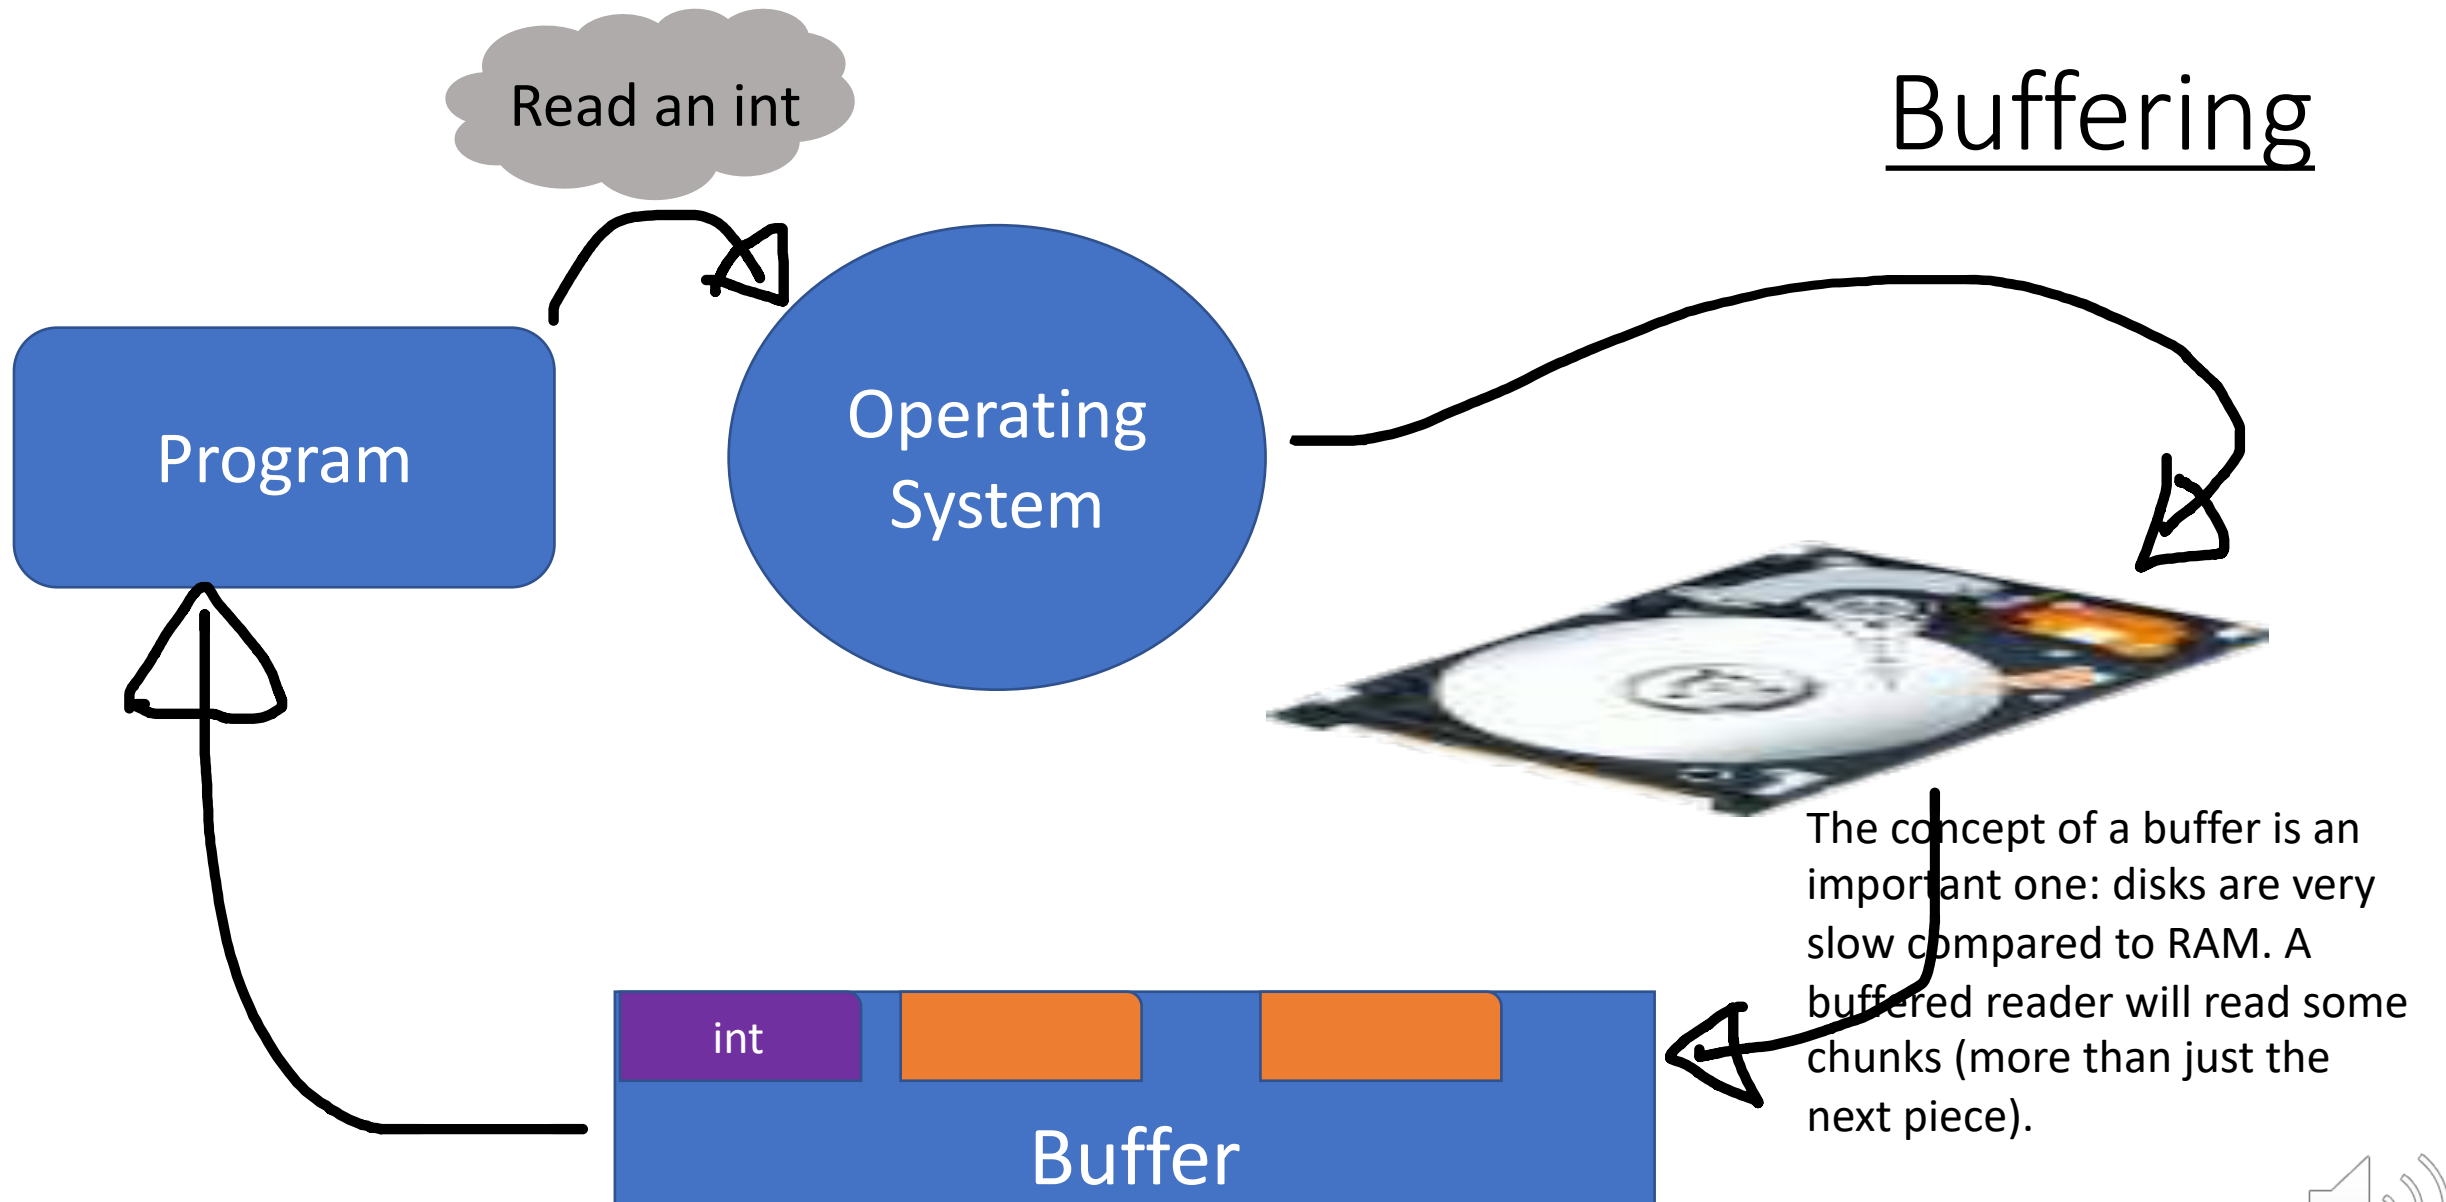
- Byte streams

Buffered or unbuffered

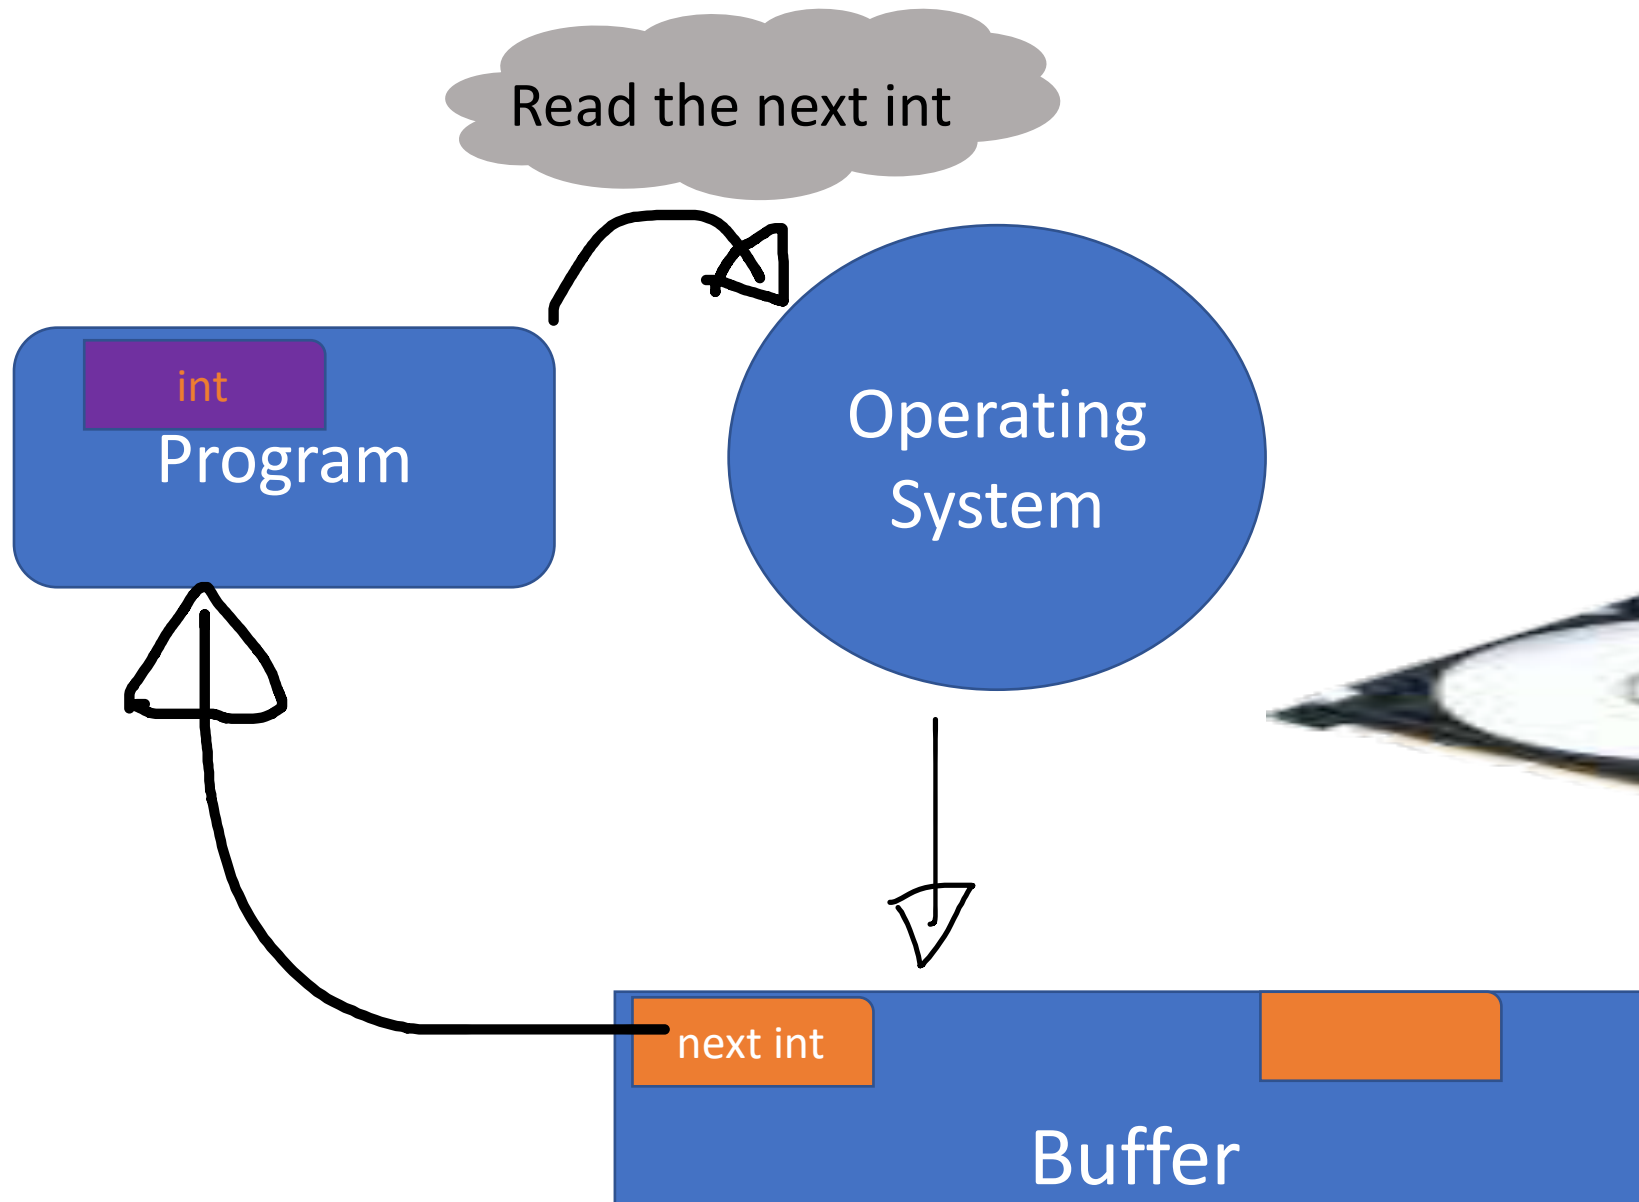
```
import java.io.*
```



# Buffering



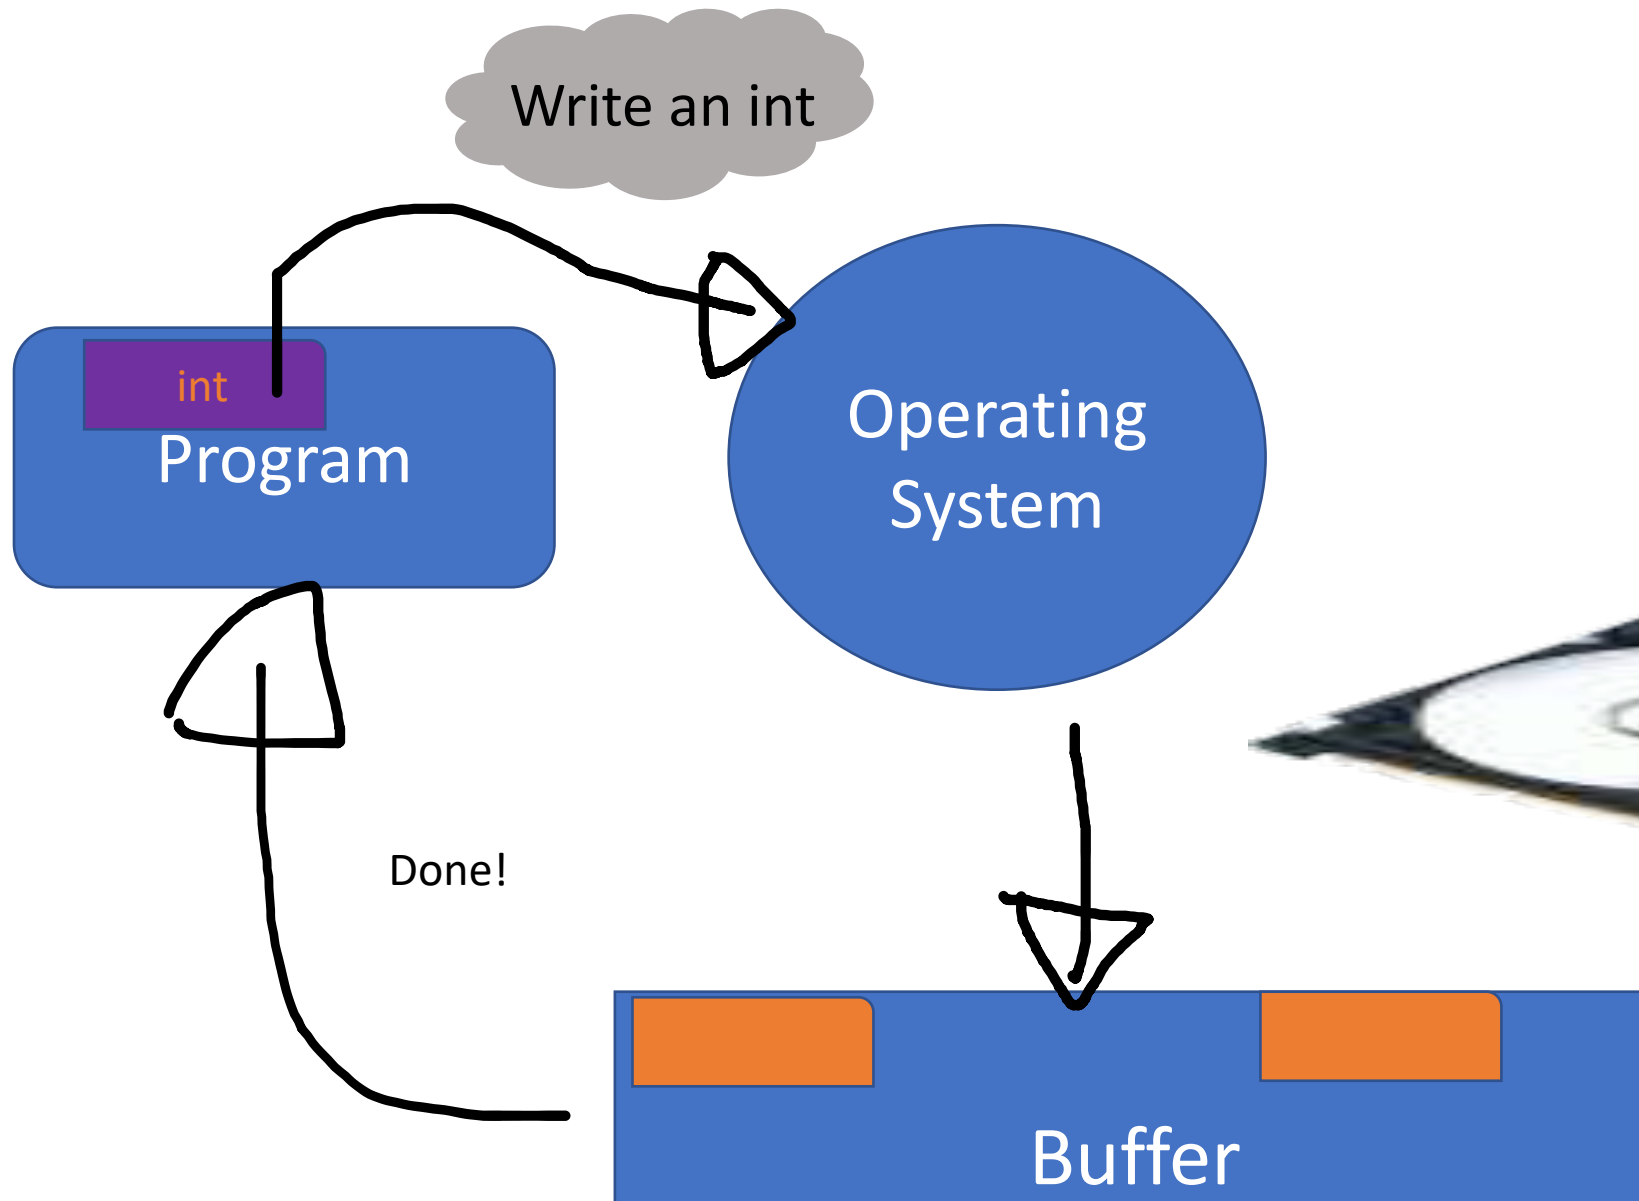
# Buffering



Then when you ask for the next int, your program runs at RAM memory speed and not at the disk speed.




# Buffering



In the same way, when you write data, it will be copied to memory, which will be much faster than writing to disk.



# Buffering



Data will be bulk transferred to the disk when the buffer is full or you close the file (or you call a method to explicitly flush the buffers).



# Buffering



Data will be bulk transferred to the disk when the buffer is full or you close the file (or you call a method to explicitly flush the buffers).



# Buffering

- What happens if the system restarts after a crash?

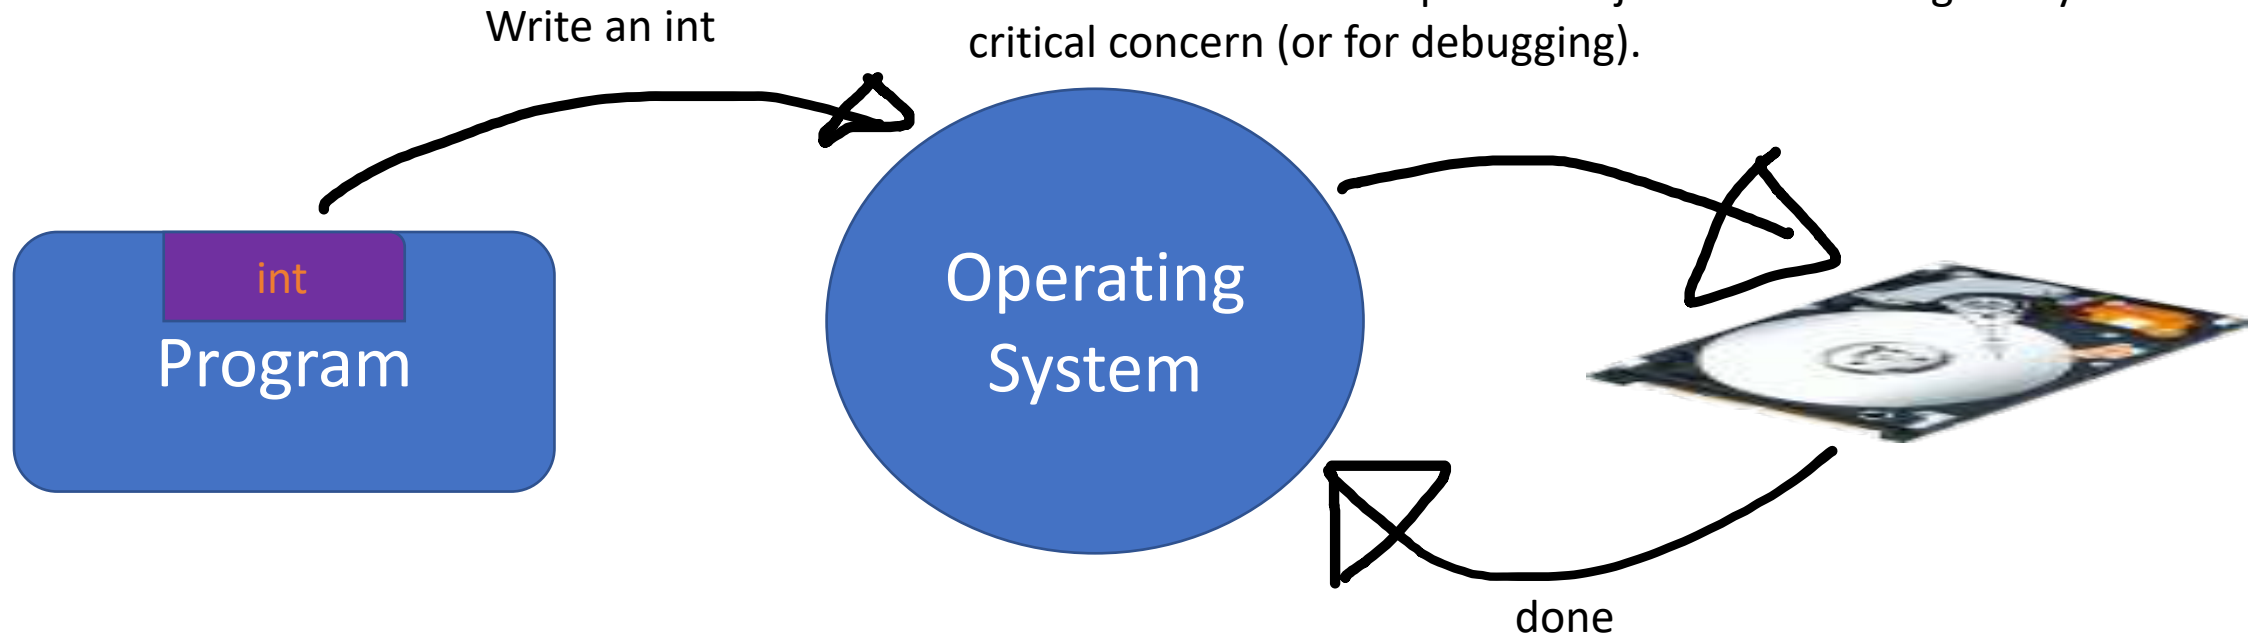
The answer of course is that what was in buffers is lost. Not a problem when reading, big problem when writing. It's not always easy to know what to replay.





# No Buffering (I/O direct to disk)

Performance: Should use buffered operations 99.9% of the time. Use unbuffered operations just when writing safely is a critical concern (or for debugging).



Instead can use slow unbuffered operations:

- Used for writing critical information
- Logs
- Messaging



# Performance - copying a 11M CSV file char by char

- Test on Mac (**internal SSD**)\*

|            |               |
|------------|---------------|
| Unbuffered | approx. 34.5s |
|------------|---------------|

|          |                   |
|----------|-------------------|
| Buffered | approx. <b>1s</b> |
|----------|-------------------|

- Test on Mac (**External USB key**)\*

|            |             |
|------------|-------------|
| Unbuffered | approx. 36s |
|------------|-------------|

|          |                     |
|----------|---------------------|
| Buffered | approx. <b>1.4s</b> |
|----------|---------------------|

- Consider if its worth it: how often does your computer crash? How bad would it be to rerun the program after restart?

\*Tests run by Stephane Faroult in previous years



# Sometimes hard to know whats going on...

- Big disk systems have their own battery protected buffers (also called a cache)

Especially with high-end storage you rarely have one level of buffering (in which case unbuffered wouldn't be what it seems). It's a bit hard sometimes to know if the data is on disk or not, and the Cloud doesn't make it any simpler.



# Unbuffered

The basis for all binary Input/Output operations are InputStreams and OutputStreams, which are unbuffered.

```
InputStream in = null;  
OutputStream out = null;
```

```
in = new InputStream(...);  
out = new OutputStream(...);
```



One thing that should not be forgotten with file operations is that it's probably the part of a program where everything can fail.

## LOTS OF THINGS CAN GO WRONG...

- Wrong file directory

- Not allowed permissions

- Content not as expected

- Hardware problem (rare)



# Unbuffered

So everything should really be done with exception handling: either with a "try with resources" or a "finally" block to make sure files are cleanly closed and not left in a "corrupted" state.

```
InputStream in = null;
OutputStream out = null;

try {
    in = new InputStream(...);
    out = new OutputStream(...);

    ...
} catch (...) {
} finally {
    if (in != null) {
        try {
            in.close();
        } catch (IOException e)
            // ignore
    }
}
```

Important: flush everything and close



# Unbuffered

```
FileInputStream in = null;  
FileOutputStream out = null;  
  
try {  
    in = new FileInputStream("filename");  
    out = new FileOutputStream("filename");  
  
    } catch (...) {  
  
    } finally {  
        ...  
    }  
}
```

Looks in the current directory unless you provide a full path

File location is always a practical problem. Use reflection and properties files.



# Buffered

To turn an unbuffered stream into a buffered one, just wrap the call to the stream constructor in a call to a buffered stream constructor.

```
BufferedInputStream in = null;
BufferedOutputStream out = null;

try {
    in = new BufferedInputStream(new InputStream(...));
    out = new BufferedOutputStream(new OutputStream(...));

} catch (...) {

} finally {
    ...
}
}
```





Next: Byte and Character Streams + Object Serialization

